

Boston House Price Prediction

March 10, 2019

```
In [2]: %matplotlib inline
import numpy as np
from matplotlib import pyplot as plt
from sklearn.datasets import load_boston
import pandas as pd
boston = load_boston()
```

```
In [3]: type(boston)
type(boston.data)
boston.feature_names
print(boston.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is used

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the

```
In [4]: bos = pd.DataFrame(boston.data)
        bos.columns = boston.feature_names
        bos['PRICE'] = boston.target
```

```
In [5]: bos.head()
```

```
Out [5]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
In [6]: bos.corr(method='spearman')
```

```
Out[6]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.571660	0.735524	0.041537	0.821465	-0.309116	0.704140	
ZN	-0.571660	1.000000	-0.642811	-0.041937	-0.634828	0.361074	-0.544423	
INDUS	0.735524	-0.642811	1.000000	0.089841	0.791189	-0.415301	0.679487	
CHAS	0.041537	-0.041937	0.089841	1.000000	0.068426	0.058813	0.067792	
NOX	0.821465	-0.634828	0.791189	0.068426	1.000000	-0.310344	0.795153	
RM	-0.309116	0.361074	-0.415301	0.058813	-0.310344	1.000000	-0.278082	
AGE	0.704140	-0.544423	0.679487	0.067792	0.795153	-0.278082	1.000000	
DIS	-0.744986	0.614627	-0.757080	-0.080248	-0.880015	0.263168	-0.801610	
RAD	0.727807	-0.278767	0.455507	0.024579	0.586429	-0.107492	0.417983	
TAX	0.729045	-0.371394	0.664361	-0.044486	0.649527	-0.271898	0.526366	
PTRATIO	0.465283	-0.448475	0.433710	-0.136065	0.391309	-0.312923	0.355384	
B	-0.360555	0.163135	-0.285840	-0.039810	-0.296662	0.053660	-0.228022	
LSTAT	0.634760	-0.490074	0.638747	-0.050575	0.636828	-0.640832	0.657071	
PRICE	-0.558891	0.438179	-0.578255	0.140612	-0.562609	0.633576	-0.547562	

	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
CRIM	-0.744986	0.727807	0.729045	0.465283	-0.360555	0.634760	-0.558891
ZN	0.614627	-0.278767	-0.371394	-0.448475	0.163135	-0.490074	0.438179
INDUS	-0.757080	0.455507	0.664361	0.433710	-0.285840	0.638747	-0.578255
CHAS	-0.080248	0.024579	-0.044486	-0.136065	-0.039810	-0.050575	0.140612
NOX	-0.880015	0.586429	0.649527	0.391309	-0.296662	0.636828	-0.562609
RM	0.263168	-0.107492	-0.271898	-0.312923	0.053660	-0.640832	0.633576
AGE	-0.801610	0.417983	0.526366	0.355384	-0.228022	0.657071	-0.547562
DIS	1.000000	-0.495806	-0.574336	-0.322041	0.249595	-0.564262	0.445857
RAD	-0.495806	1.000000	0.704876	0.318330	-0.282533	0.394322	-0.346776
TAX	-0.574336	0.704876	1.000000	0.453345	-0.329843	0.534423	-0.562411
PTRATIO	-0.322041	0.318330	0.453345	1.000000	-0.072027	0.467259	-0.555905
B	0.249595	-0.282533	-0.329843	-0.072027	1.000000	-0.210562	0.185664
LSTAT	-0.564262	0.394322	0.534423	0.467259	-0.210562	1.000000	-0.852914
PRICE	0.445857	-0.346776	-0.562411	-0.555905	0.185664	-0.852914	1.000000

```
In [7]: X = bos[['CRIM']]
        y = bos['NOX']
        type(X)
```

```
Out[7]: pandas.core.frame.DataFrame
```

```
In [8]: from sklearn.linear_model import LinearRegression
```

```
        from sklearn.model_selection import train_test_split
```


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.33, random_state=42)
#X_train, y_train = X, y
```

```
lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
Out[8]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                          normalize=False)
```

```
In [9]: plt.scatter(X, y)
```

```
Out[9]: <matplotlib.collections.PathCollection at 0x7f20ca7f07b8>
```



```
In [21]: from sklearn.metrics import r2_score
         from numpy import corrcoef
```

```
In [23]: pred = lm.predict(X_test)
         print(r2_score(X_test.values, pred))
         print(corrcoef(y_test, pred)[0])
         from scipy.stats import pearsonr
         r, p = pearsonr(X_test.values.reshape(len(X_test.values)), pred)
         r
```

```
0.98
```

```
0.933333333333
```

```
Out[23]: 0.99999999999999978
```

```
In [ ]:
```